



OBSERVABILITY

— & —

MONITORING

JUMO



PyConKe
2018



HELLO!

We are Wahome & Bildad

We are here because we love to give presentations. 😊



INTRODUCTION

What about observability & monitoring?



Reliability and stability of your services is highly dependent on understanding the state of your infrastructure and systems



Troubleshooting them effectively necessitates real-time performance, data presented at fine granularity.



METRIC, MONITORING & ALERTING INTERRELATED



Metrics, monitoring & alerting interleaved concepts that
together form the basis of a monitoring system



DEFINING METRICS, MONITORING, ALERTING

Metrics

Representing data in/from your system, a metric is an observed value of a certain quantity at a given point in time.

Types of metrics:

1. Counters
2. Gauges
3. Histograms and Timers

* Metrics != Tracing/Logs

Monitoring

The process of collecting, aggregating, & analyzing those values to improve awareness of your components' characteristics and behavior.

Two ways to collect metrics::

- X direct metric collection
- X log-based metric collection

* Monitoring != Observability

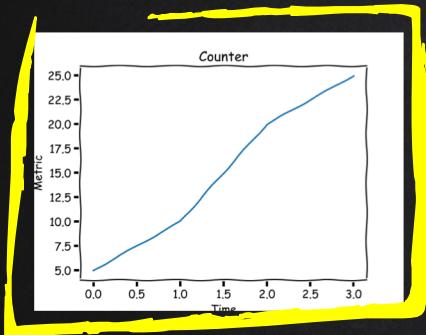
Alerting

The responsive component of a monitoring system that performs actions based on changes in metric values. Alerts definitions are composed of two components:

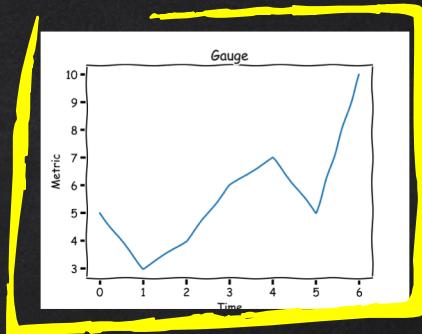
- X a metrics-based condition or threshold, and
- X an action to perform when the values fall outside of the acceptable conditions.



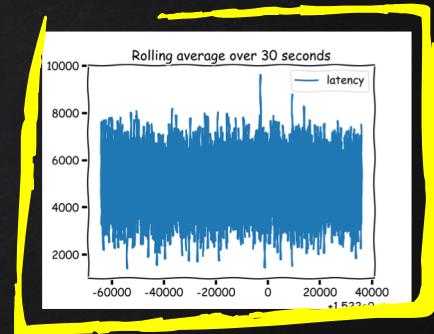
TYPES OF METRICS



A cumulative metric that represents a single monotonically increasing counter whose value can only increase or be reset to zero on restart.



A gauge is a metric that represents a single numerical value that can arbitrarily go up and down. It's a value usually has a ceiling and a floor in a certain time window.



A histogram or timer samples observations (usually things like request durations or response sizes) and counts them in configurable buckets. It also provides a sum of all observed values.



WHY DO YOU NEED A MONITORING SYSTEM?

- ✗ Visibility into the health of your systems
- ✗ Understanding normal & abnormal system & service behavior
- ✗ Capacity planning; scaling up or down
- ✗ Performance troubleshooting
- ✗ Understanding the effect of software/hardware changes
- ✗ Changing system behavior in response to a measurement
- ✗ Alerting when a system exhibits unexpected behavior



WHY HAVE YOU NOT BEEN MONITORING?

1. A false sense of security – "*A ping test will let me know if my site goes down, so all I need is a ping test!*"
2. The wrong mindset – "*This monitoring thing is a pain.*"
3. The tech debt crisis – "*We'll sort this monitoring thing once we are in production, first we ship the features!*"



MONITORING STRATEGY

What and where to monitor



WHERE TO MONITOR

1. The Host

Health or performance of an individual machine

- X CPU
- X Memory
- X Disk space
- X Processes

2. The Application

Indicators of the health, performance, or load of an application

- X Error and success rates
- X Service failures & restarts
- X Response latencies
- X Resource usage

3. The Server Pool

System-wide resource usage data

- X Pooled resource usage
- X Scaling adjustment indicators
- X Degraded instances

4. Network & Connectivity

Accessibility to other machines

- X Connectivity
- X Error rates and packet loss
- X Latency
- X Bandwidth utilization

5. External Dependencies

E.g. DNS, SSL Certs

- X Service status and availability
- X Success and error rates
- X Run rate and operational costs
- X Resource exhaustion

6. The User Experience

The last but certainly not least layer in the stack is the user.

At this layer, we're interested in the behavior of the full stack as a user would experience it e.g. Apdex score



WHAT TO MONITOR



Monitor Potential Bad Things

Alert before they happen



Monitor Actual Bad Things

Alert when they do happen, which is, unfortunately, inevitable



Monitor Good Things

Alert when they stop happening



Tuning & Continuous Improvement

Monitoring is an iterative process – you're not expected to get it all right the very first time around.



QUALITIES OF A GOOD MONITORING SYSTEM

- X Independent from most other infrastructure
- X Reliable & Trustworthy
- X Summary & detail views; easy to use
- X Maintaining historical data; patterns
- X Aggregation; correlate factors from different sources
- X Integration; easy to start tracking new metrics or infrastructure
- X Flexible & powerful alerting



DESIGNING FOR MONITORING

Monitoring ought to be a first-class citizen. Instrumentation should be an integral part of your code.



Metrics instrumentation & tracing/logs will
more often than not occur at the boundary
of your components where interfaces are
exposed



Every piece of knowledge or logic
ought to have a single,
unambiguous representation.

Write code once, use it often.
Change code in one place, see the
change in all instances.



HEALTH CHECK ENDPOINT



A URL endpoint which monitoring systems & load balancers can use to determine application health. If the application returns a success status it is considered healthy.



LET'S WALK THE TALK

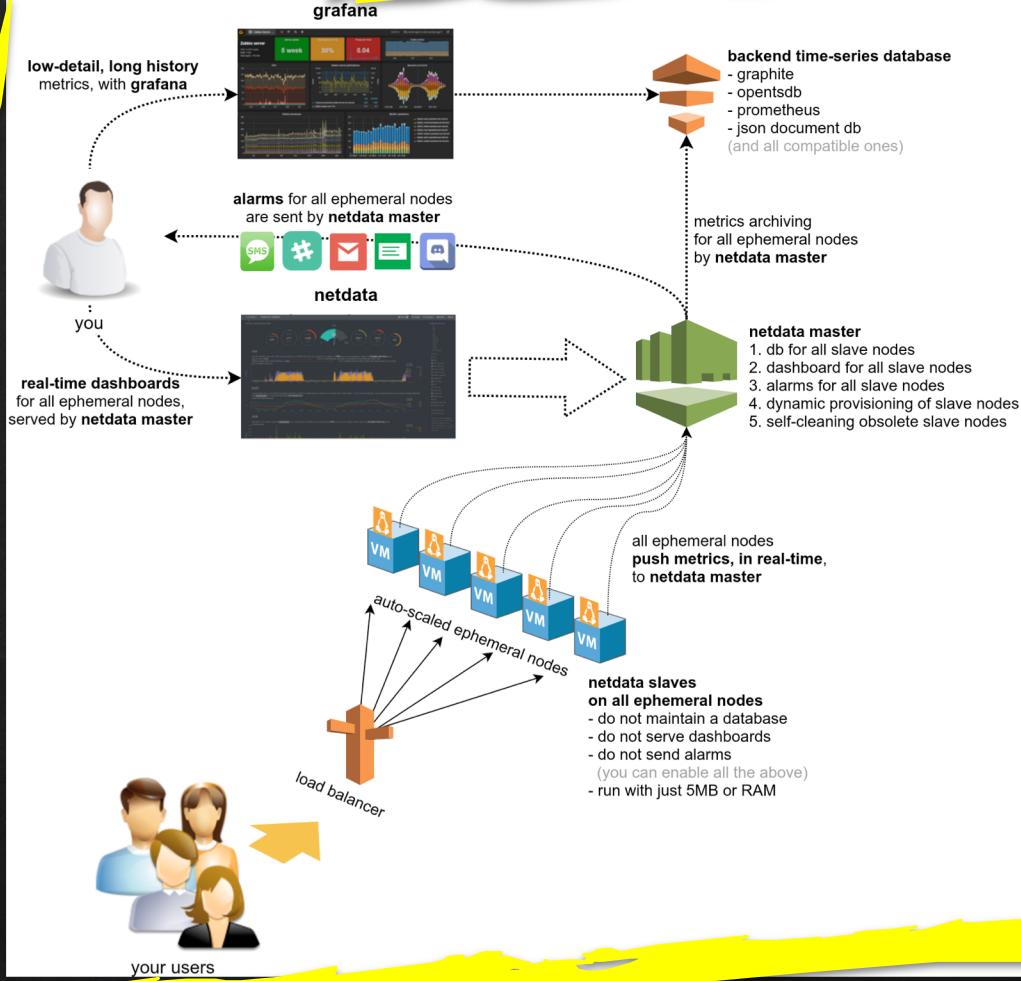
Demo Time



MESSAGING/NOTIFICATION SERVICE

To adequately demonstrate implementation of a monitoring system in a Python application/service, we have built a [messaging service](#) and integrated [netdata](#) into the various components (nodes) for metrics instrumentation.

All metrics data from the nodes is streamed into the netdata master from which it's visualized (real-time analysis on netdata) & streamed to [prometheus](#) for storage (and further visualization with [grafana](#))





WHERE TO GET THE CODE

Link: <https://github.com/kwahome/pycon-monitoring-workshop>

You'll need docker & compose.
Setup details in the README



REFERENCES

- X [An introduction to metrics, monitoring and alerting by DigitalOcean](#)
- X [Zen and the art of system monitoring](#)
- X [Metrics monitoring and python](#)
- X [Grafana](#)
- X [Prometheus](#)
- X [Netdata](#)



THANKS!

Any questions?

@kwahome_

@bnamawa